

Hands-On: Kubernetes Dashboard

Step 1:

Launch 2 instances with the following configuration: ubuntu 20.04 ami, t2.medium, sg: all traffic. ubuntu 20.04 ami, t2.micro , sg: all traffic
To Install Kubernetes use the following commands:

On Master and Worker node:

```
sudo su
apt-get update
apt-get install docker.io -y
service docker restart
curl -s https://packages.cloud.google.com/apt/doc/apt-key.gpg | apt-key
add -
echo "deb http://apt.kubernetes.io/ kubernetes-xenial main"
>/etc/apt/sources.list.d/kubernetes.list
apt-get update
apt install kubeadm=1.20.0-00 kubectl=1.20.0-00 kubelet=1.20.0-00 -y
```

Step 2: On both master and worker nodes run the above command:

- 2.1. sudo su
- 2.2. create a script file kubernetes.sh
- 2.3. to execute the script file: bash kubernetes.sh

On Master:

Step 3:

Creating cluster:

Initializing kubeadm on master using:

```
kubeadm init --pod-network-cidr=192.168.0.0/16
```

```
[bootstrap-token] configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubectl apply -f [podnetwork].yaml" with one of the options listed at:
https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.18.147:6443 --token rb42wf.rli4srhyw2cark8s \
--discovery-token-ca-cert-hash sha256:bc23c6c2902c4c5f514ad86075b474e9c5b7060cd42e275499297c959aa8bd41
root@ip-172-31-18-147:/home/ubuntu#
```

Copy the token and paste it into the worker node

```
root@ip-172-31-13-193:/home/ubuntu# kubeadm join 172.31.18.147:6443 --token rb42wf.rli4srhyw2cark8s \
--discovery-token-ca-cert-hash sha256:bc23c6c2902c4c5f514ad86075b474e9c5b7060cd42e275499297c959aa8bd41
[preflight] Running pre-flight checks
[WARNING IsDockerSystemdCheck]: detected "cgroupfs" as the Docker cgroup driver. The recommended driver is "systemd". Please follow the guide at https://kubernetes.io/docs/setup/cni/
[WARNING SystemVerification]: this Docker version is not on the list of validated versions: 20.10.7. Latest validated version: 19.03
[preflight] Reading configuration from the cluster...
[preflight] FYI: You can look at this config file with 'kubectl -n kube-system get cm kubeadm-config -o yaml'
[kubelet-start] Writing kubelet configuration to file "/var/lib/kubelet/config.yaml"
[kubelet-start] Writing kubelet environment file with flags to file "/var/lib/kubelet/kubeadm-flags.env"
[kubelet-start] Starting the kubelet
[kubelet-start] Waiting for the kubelet to perform the TLS Bootstrap...

This node has joined the cluster:
* Certificate signing request was sent to apiservert and a response was received.
* The Kubelet was informed of the new secure connection details.

Run 'kubectl get nodes' on the control-plane to see this node join the cluster.

root@ip-172-31-13-193:/home/ubuntu#
```

Step 4:

On Master:

We need to run the below commands:

```
exit
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

Note: In case we want to retrieve the join token use the below-mentioned command.

```
kubeadm token create --print-join-command
```

Step 5:

On Master:

```
kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml  
kubectl apply -f  
https://raw.githubusercontent.com/kubernetes/ingress-nginx/controller-v0.  
49.0/deploy/static/provider/baremetal/deploy.yaml
```

To list all nodes :

```
kubectl get nodes
```

Our Kubernetes installation and configuration is complete.

It shows the nodes but the status is not ready
because we have not installed the network plugin.

To install network plugin, run the below commands:

```
kubectl apply -f https://docs.projectcalico.org/manifests/calico.yaml  
kubectl apply -f https://raw.githubusercontent.com/kubernetes/ingress-  
nginx/controller-v0.49.0/deploy/static/provider/baremetal/deploy.yaml
```

Check the nodes for its state after installing network plugins.

```
kubectl get nodes
```

Thus, we have successfully installed Kubernetes.

Step 6:


Run the below command to create a dashboard:

```
kubectl apply -f  
https://raw.githubusercontent.com/kubernetes/dashboard/v2.0.0/aio/deploy/recommen  
ded.yaml
```

Then edit the service:

```
kubectl edit service kubernetes-dashboard -n kubernetes-dashboard
```

```
Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: v1
kind: Service
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"Service","metadata":{"annotations":{},"labels":{"k8s-app":
s-dashboard"},"spec":{"ports":[{"port":443,"targetPort":8443}],"selector":{"k8s-app":"kuber
creationTimestamp: "2022-01-13T07:16:45Z"
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
  namespace: kubernetes-dashboard
  resourceVersion: "1539"
  uid: 01f6f13a-d1bf-45b0-afc4-ca9713168cac
spec:
  clusterIP: 10.105.231.218
  clusterIPs:
  - 10.105.231.218
  ipFamilies:
  - IPv4
  ipFamilyPolicy: SingleStack
  ports:
  - port: 443
    protocol: TCP
    targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard
  sessionAffinity: None
  type: ClusterIP
status:
```



A diagram illustrating the modification of the service type. A red box highlights the 'ClusterIP' value under the 'type' field in the 'spec' section. A blue arrow points from this box to another red box on the right containing the text 'Change ClusterIP to NodePort'.

Note: We need to change the type from ClusterIP to NodePort. The below given image contains the modified service file.


```

apiVersion: v1
kind: Service
metadata:
  annotations:
    kubectl.kubernetes.io/last-applied-configuration: |
      {"apiVersion":"v1","kind":"Service","metadata":{"annotations":{},"labels":{"k8s-app":"kubernetes-dashboard"},"spec":{"ports":[{"port":443,"targetPort":8443}],"selector":{"k8s-app":"kubernetes-dashboard"},"status":{"loadBalancer":{}}}}
  creationTimestamp: "2022-01-13T07:16:45Z"
  labels:
    k8s-app: kubernetes-dashboard
  name: kubernetes-dashboard
  namespace: kubernetes-dashboard
  resourceVersion: "1997"
  uid: 01f6f13a-d1bf-45b0-afc4-ca9713168cac
spec:
  clusterIP: 10.105.231.218
  clusterIPs:
    - 10.105.231.218
  externalTrafficPolicy: Cluster
  ipFamilies:
    - IPv4
  ipFamilyPolicy: SingleStack
  ports:
    - nodePort: 31707
      port: 443
      protocol: TCP
      targetPort: 8443
  selector:
    k8s-app: kubernetes-dashboard
  sessionAffinity: None
  type: NodePort
status:
  loadBalancer: {}

```

Note: The editor used is vim. Therefore to save and exit we need to press ESC and then: wq

```

root@ip-172-31-23-197:/home/ubuntu# kubectl edit service kubernetes-dashboard -n kubernetes-dashboard
service/kubernetes-dashboard edited

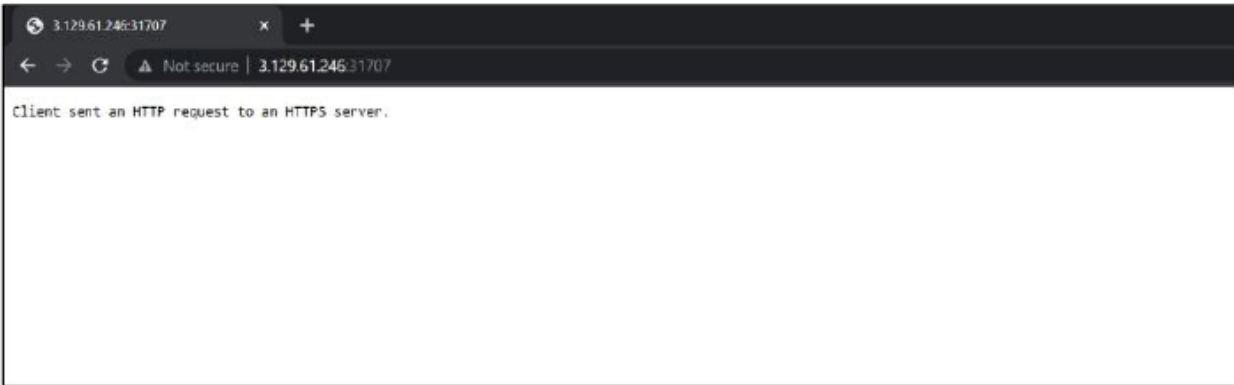
```

kubectl get svc -n kubernetes-dashboard

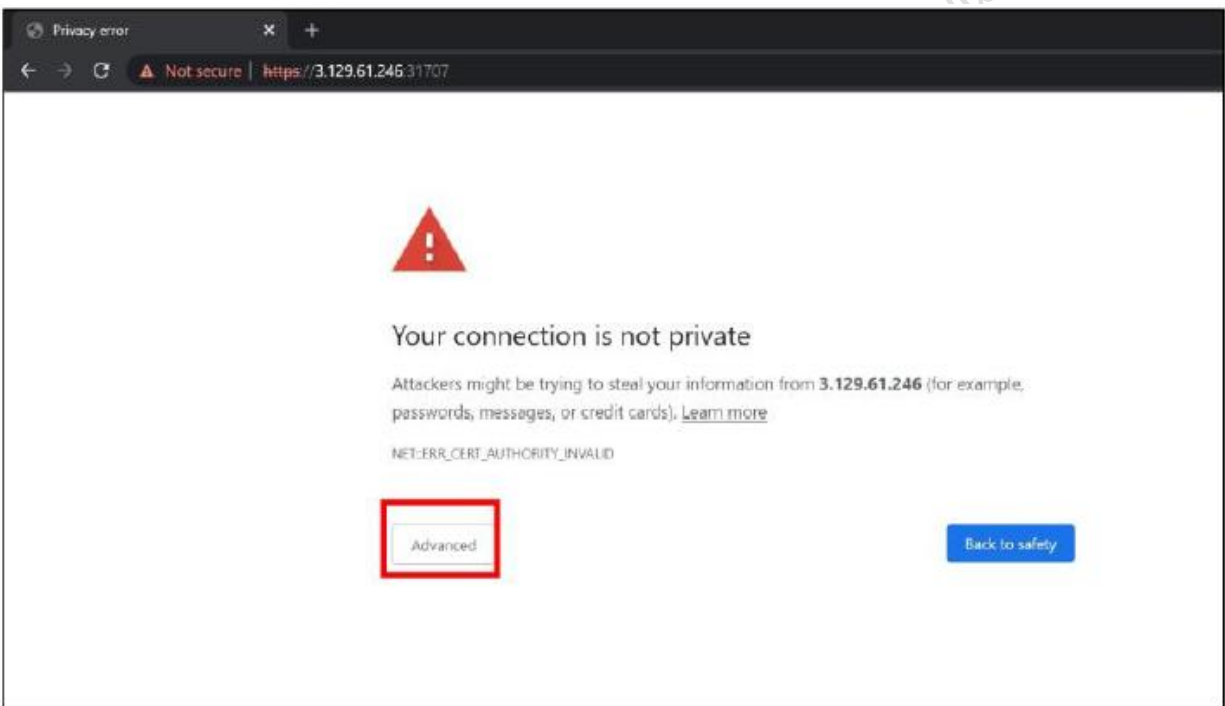
Step 7:

Now go to a browser and paste the ip along with the port

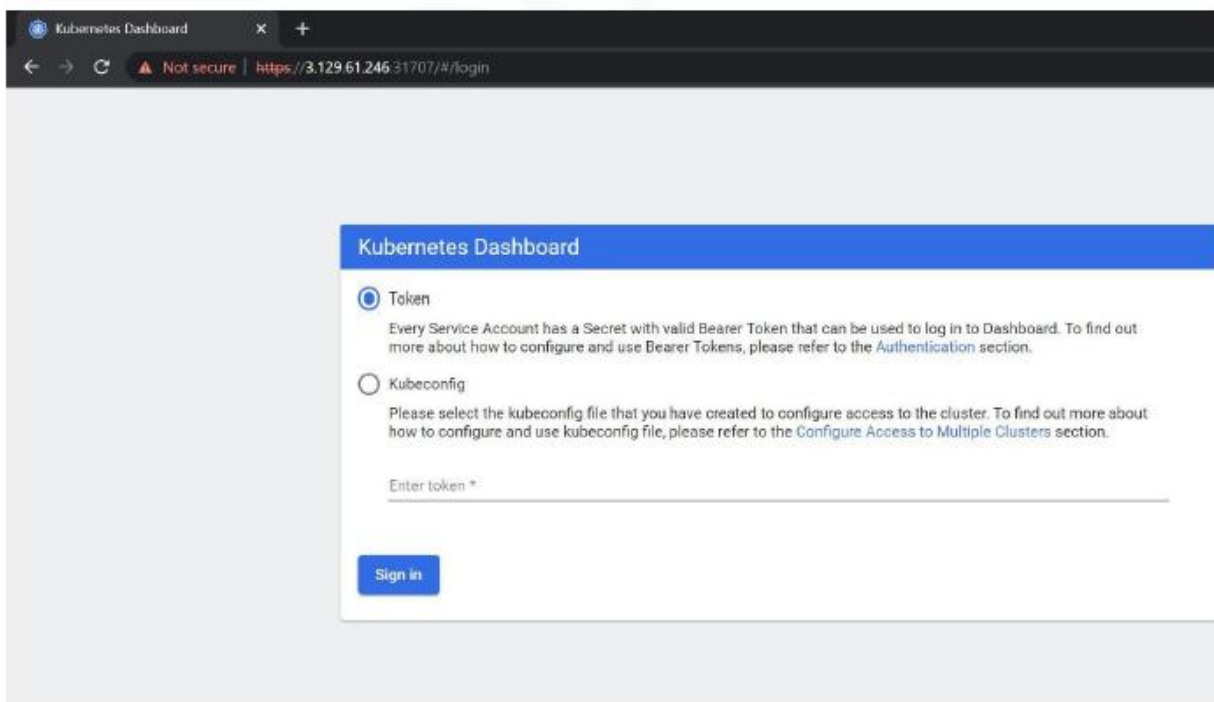
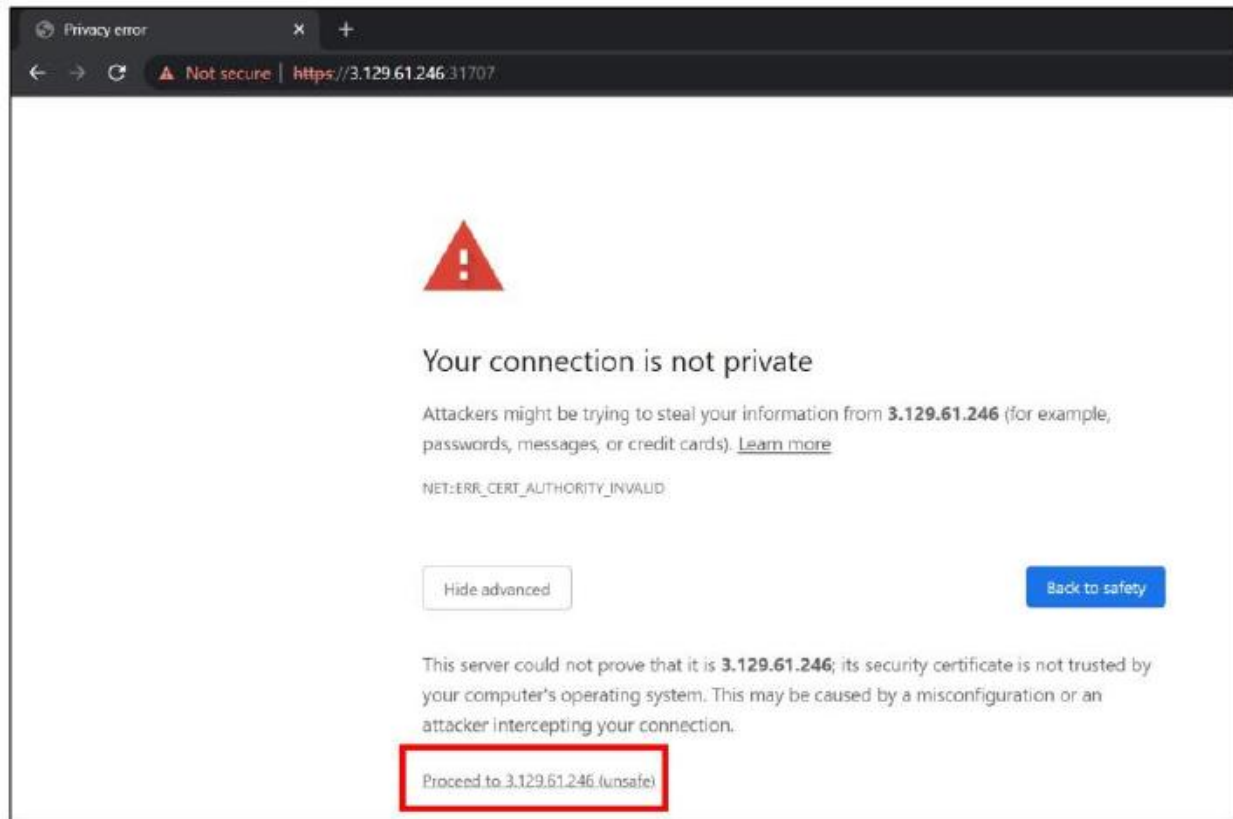
<https://<ip-of-master>>



We need to click on Advance option to access the webpage



Click on the Proceed option to open the dashboard



Copy the token and paste in the Kubernetes dashboard.

Kubernetes Dashboard

☒ Token

Every Service Account has a Secret with valid Bearer Token that can be used to log in to Dashboard. To find out more about how to configure and use Bearer Tokens, please refer to the [Authentication](#) section.

☐ Kubeconfig

Please select the kubeconfig file that you have created to configure access to the cluster. To find out more about how to configure and use kubeconfig file, please refer to the [Configure Access to Multiple Clusters](#) section.

Enter token *

.....

[Sign in](#)

Then click on Sign in. The dashboard is created.

The screenshot shows the Kubernetes Dashboard Overview page in a web browser. The browser address bar shows the URL `https://3.129.61.246:31707/#/overview?namespace=default`. The dashboard has a blue header with the 'kubernetes' logo and a search bar. A left sidebar contains navigation links: Overview, Cluster Roles, Namespaces, Nodes, Persistent Volumes, Storage Classes, Namespace (set to 'default'), Workloads, Cron Jobs, Daemon Sets, and Deployments. The main content area is divided into two sections: 'Discovery and Load Balancing' and 'Config and Storage'. The 'Discovery and Load Balancing' section contains a 'Services' table with one entry: 'kubernetes' in the 'default' namespace, with labels 'component: apiserver' and 'provider: kubernetes', and a cluster IP of '10.96.0.1'. The 'Config and Storage' section contains a 'Config Maps' table with one entry: 'kube-root-ca.crt' in the 'default' namespace.

Name	Namespace	Labels	Cluster IP	Internal Endpoints	External Endpoints	Created
kubernetes	default	component: apiserver provider: kubernetes	10.96.0.1	kubernetes:443 TCP kubernetes:0 TCP	-	33 minutes ago

Name	Namespace	Labels	Created
kube-root-ca.crt	default	-	33 minutes ago